

Deep reinforcement learning-based optimal deployment of IoT machine learning jobs in fog computing architecture

Omid Bushehrian, Amir Moazeni

By: Mohammad Kadkhodaei

1404-04-07, rev 1.0

Abstract

- Many IoT applications are **machine learning jobs** that collect and analyze sensor measurements.
- These machine learning jobs consist of **distributed tasks** that work collaboratively to build models in a **federated manner**.
- The problem of determining the optimal number and the coverage of distributed tasks of an IoT machine learning job has not been addressed previously
- This paper proposes a **two-phased method** for adaptive task **creation** and **deployment** of IoT ML jobs over a heterogeneous multi-layer fog computing architecture (adaptively and efficiently).

Abstract

- In the **1st phase**, the optimal number of tasks and their respective sensor coverage is determined using a Deep Reinforcement Learning (DRL) > **Task Creation**, using > **Deep Deterministic Policy Gradient (DDPG)**
- In the **2nd phase**, the tasks are deployed over the heterogeneous multi-layer fog computing architecture using a **greedy** deployment method > **Task Deployment**
- **DDPG Sensor positioning Component + Task creation Component + Greedy Deployment Component**
- The task creation and deployment is formulated as a **3- optimization problem**:
 - 1) min the deployment latency
 - 2) min the deployment cost
 - 3) min the evaluation loss

Result

- The proposed two-phased DRL-based method could outperform the **Edge-IoT** and **Cloud-IoT** baseline methods by improving the **total deployment score** up to **32%**.

$$\text{Maximize } \sum_{j_i \in J} deploymentScore_j$$

$$deploymentScore_j = -1 \times (w_1 \times latencyFactor_j^N + w_2 \times costFactor_j^N + w_3 \times lossFactor_j^N)$$

- Simulation: Python + Keras/TensorFlow, COREi7, 32MB RAM

Table 1 The comparison of the related works using different criteria. Time, cost, accuracy, resource utilization, and energy consumption are denoted by T, C, A, R, and E respectively in the objectives column

Paper	Category	FL-based	Application architecture	Task creation	Infrastructure	Algorithm	Objectives				
							T	C	A	R	E
[3]	Mathematical/(meta-) heuristic methods	No	Microservice	No	Mobile edge	Branch-and- bound	×	●	×	×	×
[20]		No	DAG	No	Multi-layer Fog	Dynamic scheduling	●	×	×	×	×
[21]		No	Microservice	No	Fog	PSO	●	●	×	×	×
[22]		No	Independent multi-task	No	Fog	AEO	●	×	×	×	×
[23]		No	Independent multi-task	No	Fog	GWO	●	●	×	×	×
[24]		No	Independent services	No	Fog	WOA	●	×	×	×	●
[25]		No	Worker/parameter server	Yes(non-adaptive)	Cluster	Approximation algorithms	●	×	×	×	×
[4]	DRL-based	Yes	FL- three-layer Architecture	Yes (fixed task number)	Edge/Cloud	MDRL	●	×	●	×	●
[10]		No	DAG	No	Heterogeneous fog	DDRL	●	×	×	×	●
[26]		No	DAG	No	Fog	RL	●	●	×	×	×
[27]		No	Independent multi-task	No	Mobile edge	DDLO	●	×	×	●	●
[28]		No	Independent multi-task	No	Edge	DRL	●	×	×	●	×
[29]		No	Independent multi-task	No	Edge	DQN	●	×	×	●	×
[30]		No	Microservice	No	Heterogeneous fog	A3C	●	●	×	×	×
[31]		No	Cloud jobs	No	Cloud	Actor-Critic DRL	●	×	×	×	●
[32]		No	PoW single task	No	MCS	DRL	●	×	×	×	●
[33]		No	C-PoW single task	No	MCS	DRL	●	×	×	×	●
[34]		No	Single task	No	MEC	DRL	×	●	×	●	×
Proposed method	Ensemble	Yes	FL- two-layer architecture	Yes (adaptive)	Multi-layer heterogeneous Fog	DDPG-Greedy	●	●	●	×	×

$$\textit{Maximize} \sum_{j_j \in J} \textit{deploymentScore}_j$$

$$\textit{deploymentScore}_j = -1 \times (w_1 \times \textit{latencyFactor}_j^N + w_2 \times \textit{costFactor}_j^N + w_3 \times \textit{lossFactor}_j^N)$$

$$\textit{latencyFactor}_j = \sum_{h_k \in F} \sum_{s_i \in M_j} X_{i,j,k} \times L_k$$

$$\textit{costFactor}_j = \textit{compCost}_j + \textit{commCost}_j$$

$$\textit{lossFactor}_j = \left| T_j \right|$$

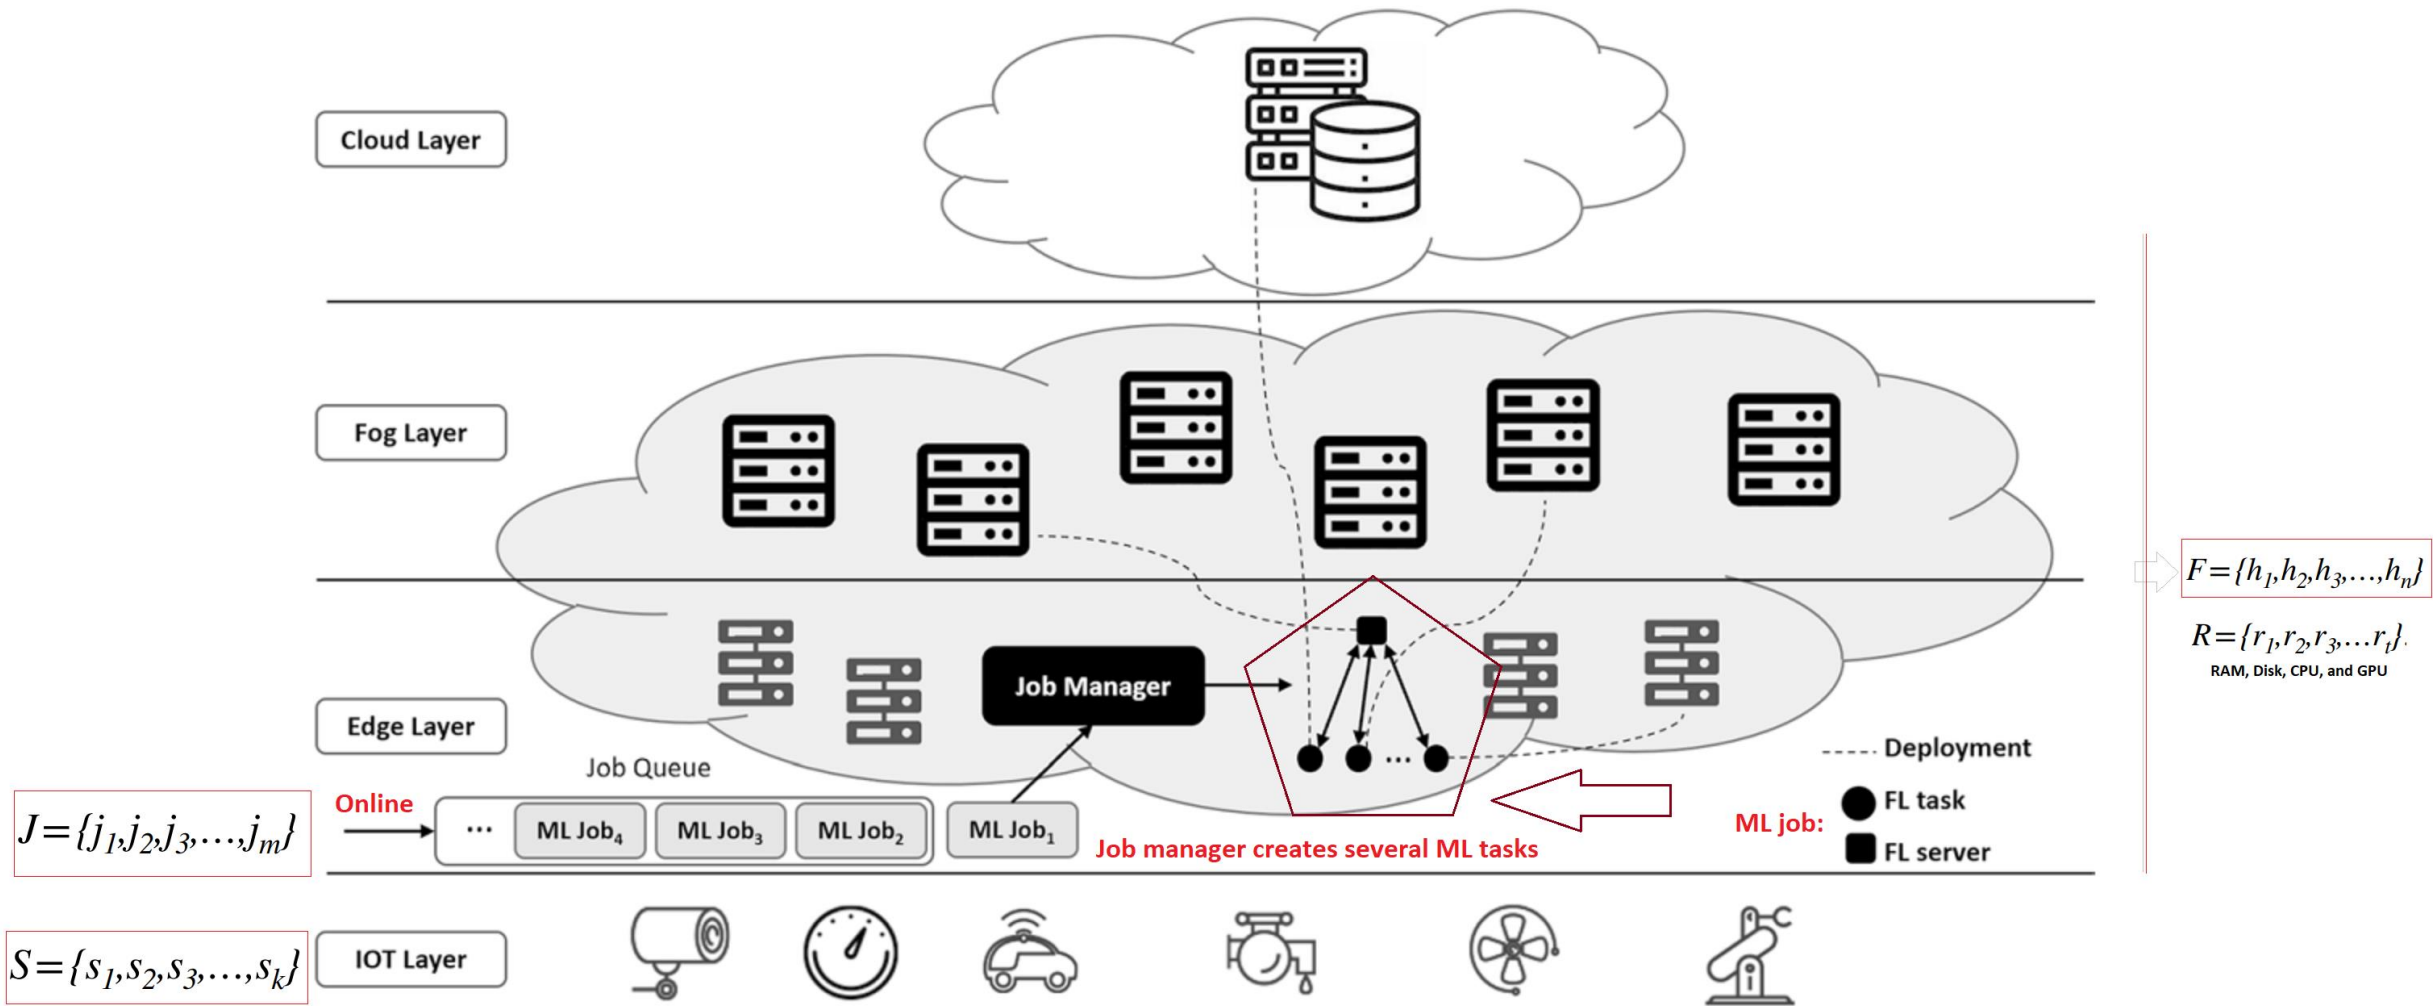
Jobs, Job Manager, Tasks set

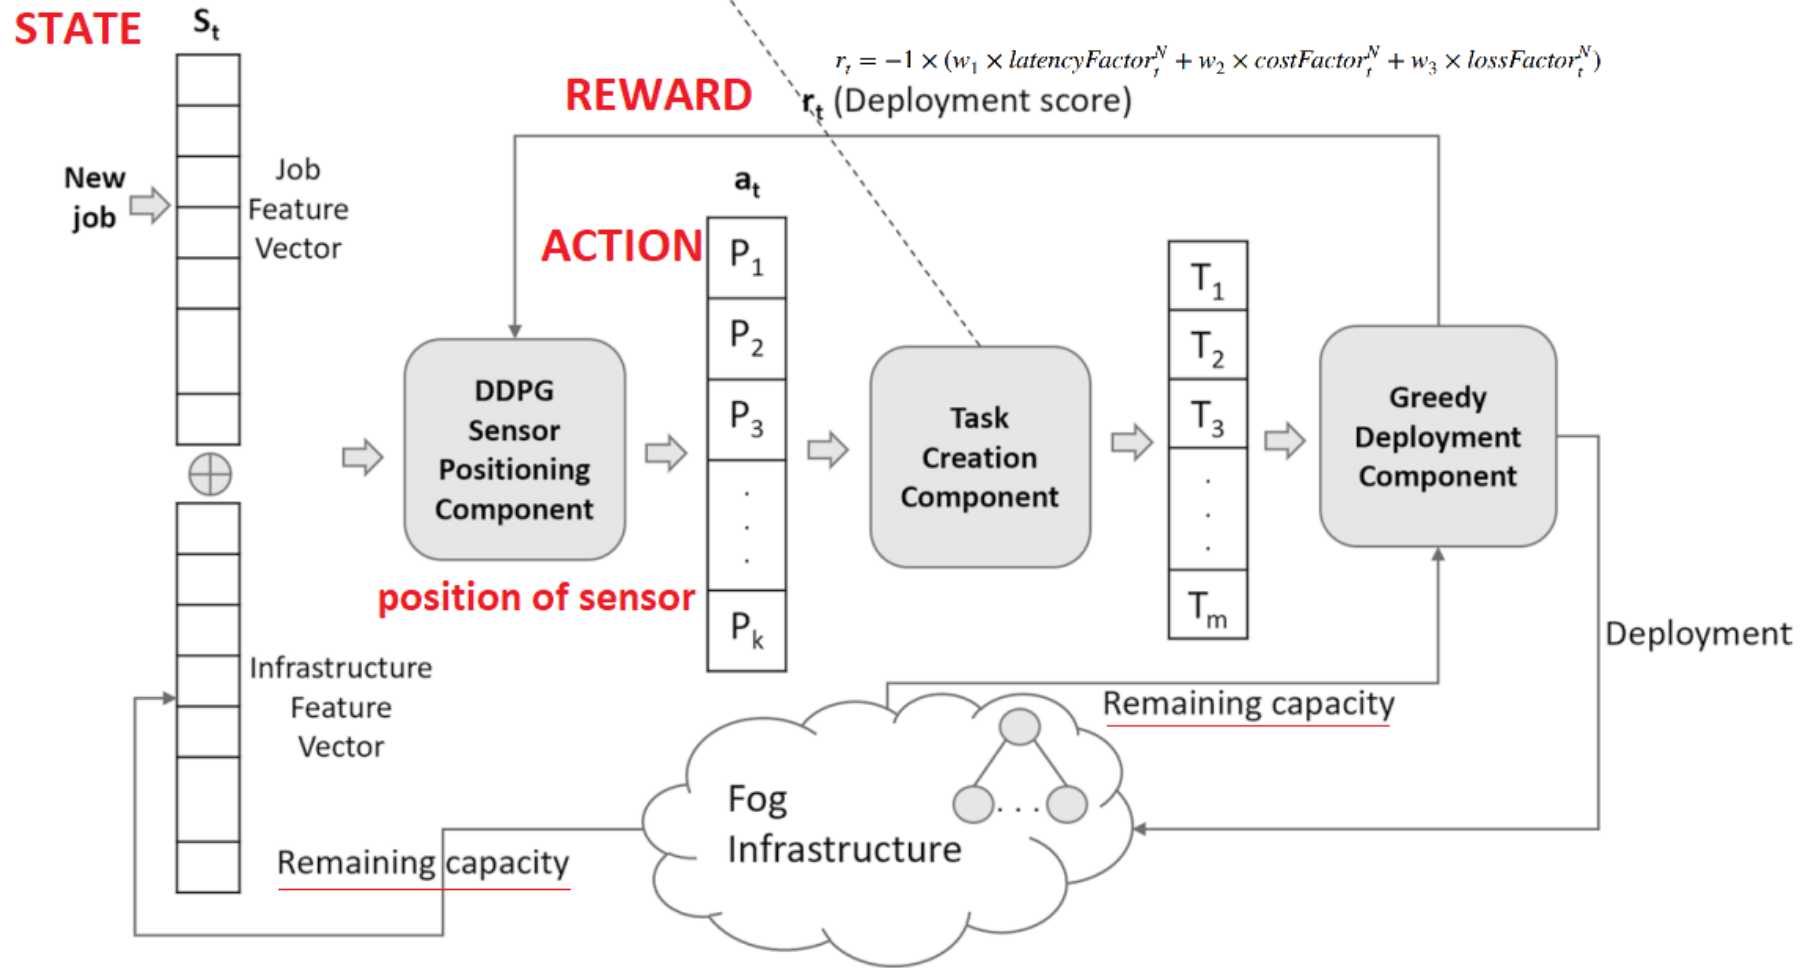

- Both offline and **online** deployment problem are NP-Hard
- Jobs are received by the **job manager** over time and the job manager makes the deployment once a new job is received.

IoT ML jobs, $J = \{j_1, j_2, j_3, \dots, j_m\}$

the aim is to define the


task set $T_j = \left\{ t_k^{n,j} \mid h_k \in F \text{ and } 1 \leq n \leq g_j \right\}$ for job j_j





DRL

- 1) Value-based methods : Q-learning, SARSA, DQN, and company
- 2) Policy-based methods : REINFORCE, ...
- 3) Actor-Critic, learn both policies and value functions : A3C, DDPG, ...



DDPG : Deep Deterministic Policy Gradient

- In 2014, Silver et al. released the deterministic policy gradient (DPG) algorithm
- DDPG was introduced in 2015 in a paper titled “Continuous control with deep reinforcement learning.” The paper was authored by Timothy Lillicrap (et al.) while he was working at Google DeepMind as a research scientist
- Is a popular **model-free**, **off-policy** and **actor–critic** DRL algorithm
- It consists of **4- DNNs**: actor, critic, target actor, target critic

Greedy Algorithm

- Iteratively places tasks on the nodes from the large to the small sizes
- Size = total computation demand
- the FL server is deployed first.
- A task deployment score is computed based on the latency, computation and communication costs
- The node with the best deployment score is selected for the task placement

The greedy job deployment algorithm

Data: T : FL Tasks, F : Infrastructure nodes

Result: deployment: Task's deployment

```
1 foreach  $task \in T$  do
2   |  $demand[task] \leftarrow calcDemand(task);$ 
3 end
4 while  $T$  contains undeployed tasks do
5   | if  $FLserver$  is not deployed then
6     |  $selectedTask \leftarrow FLserver;$  First FLserver
7   | else
8     |  $selectedTask \leftarrow nextLargestTask(demand);$  nextLargestTask
9   | end
10  | foreach  $h \in F$  do
11    | if  $demand(selectedTask) \leq remainingCapacity(h)$  then
12      |  $score[h] \leftarrow calcScore(h, selectedTask);$ 
13    | else Check Have Capacity
14      |  $score[h] \leftarrow -\infty;$ 
15    | end
16  | end
17  |  $selectedNode \leftarrow nodeWithMaxScore(score);$ 
18  |  $deployment[selectedTask] \leftarrow selectedNode;$ 
19  | Update  $selectedNode$  capacity
20  | Mark  $selectedTask$  as deployed
21 end
```

Federated Learning

- Locally-trained model parameters are received and aggregated by a centralized server to update the global model parameters using a model aggregation algorithm like : Federated averaging (**FedAvg**)
- Using : Flower framework + (LSTM) neural network in Python
- Dataset: air-quality measurements (7-features)

Future works

1. Replacing the greedy algorithm with a second DRL-based deployment algorithm
2. Support unsupervised ML jobs

My Idea

- Improve Queue structure
- Determine Jobs structure more clearly
- Add priority to jobs
- Use Meta Heuristics instead of Greedy algorithm
- Use another DNN instead of Greedy algorithm
- Redundant job manager + Fault tolerance
- Consider energy consumption